

Mosaic: Generating The New Yorker Style Cartoons using Text-to-Image Diffusion Models

Mehul Sudrik, Rajesh Nagula, Utsav Oza

ECE-GY 7123 Deep Learning

Problem Statement

“Sometimes I wonder if I’m too old to be a cartoonist, but then I remember that I’m just not funny enough.”

This self-deprecating joke highlights the challenges of creating The New Yorker style cartoons, that uniquely combines a whimsical art style, witty humor, and a subtle commentary on modern life. Cartoon enthusiasts and creative professionals alike know that creating such cartoons can be a daunting task that requires both artistic talent and a knack for satire. In this project, we aim to explore techniques to simplify the cartoon creation process by using Text-to-Image Diffusion models to specifically generate high-quality The New Yorker style cartoons from natural language captions.

Literature Survey

Text-to-image diffusion models are a type of generative model that can produce high-quality images from textual descriptions. Diffusion (Rombach et al. 2020) is a probabilistic process that involves gradually adding noise to an image until it becomes completely random, and then gradually removing the noise until it becomes the target image. The goal of diffusion models is to learn the latent structure of a dataset by modeling the way in which data points diffuse through the latent space.

The challenge, however, is that it is unclear how the diffusion process can be directly exercised to generate images of specific unique concepts, or compose them in new roles and novel scenes. Textual Inversion (Gal et al. 2022) allows us to teach text-to-image diffusion models new concepts - it takes a small number of images of a user-provided concept, like an object or a style, and learns to represent it through new “words” in the embedding space of a frozen text-to-image model. These “words” can be composed into natural language sentences, guiding personalized creation in an intuitive way.

Project Considerations

Dataset

A key challenge in fine-tuning a text-to-image diffusion model is the requirement of a large and diverse dataset, which can be difficult to obtain for a specific domain, especially if it involves rare or complex visual concepts. Textual Inversion will allow us to teach an image generator a

specific visual concept using very few image examples. As such, we’ll be relying primarily on the dataset derived for The New Yorker’s Cartoon Captioning Contest (Hessel et al. 2022), which is composed of raw cartoon images that are mapped to quality-based ranked caption choices and an explanation describing the underlying humor of the cartoon.

Model

In this project, we will be utilizing Stability.ai’s implementation of text-to-image diffusion model called Stable Diffusion. Specifically, we aim to explore the application of a fine-tuning technique called Textual Inversion on Stable Diffusion. This will involve creating a new text encoder and training it on our new embeddings, i.e. the captions for the cartoons we’d like to generate. The training algorithm will effectively take a sample from the output of the frozen Stable Diffusion image encoder’s latent distribution for a training image, add noise to that sample, and then pass that noisy sample to the frozen diffusion model. Our final goal state will be when the model is able to separate the noise from the sample using the text encoding as hidden state.

Evaluation

Considering the goal of this project is to explore techniques to simplify the cartoon creation process using Stable Diffusion, evaluating the generated outputs as such will be subjective in nature. We aim to use a combination of both qualitative (human assessment) and quantitative metrics (CLIP scores) to gauge the effectiveness of the model.

References

- Gal, R.; Alaluf, Y.; Atzmon, Y.; Patashnik, O.; Bermano, A. H.; Chechik, G.; and Cohen-Or, D. 2022. An Image is Worth One Word: Personalizing Text-to-Image Generation using Textual Inversion. *arXiv:2208.01618*.
- Hessel, J.; Marasović, A.; Hwang, J. D.; Lee, L.; Da, J.; Zellers, R.; Mankoff, R.; and Choi, Y. 2022. Do Androids Laugh at Electric Sheep? Humor “Understanding” Benchmarks from The New Yorker Caption Contest. *arXiv preprint arXiv:2209.06293*.
- Rombach, R.; Blattmann, A.; Lorenz, D.; Esser, P.; and Ommer, B. 2020. High-Resolution Image Synthesis with Latent Diffusion Models. *arXiv:2112.10752*.